

COMPTE RENDU JAVA PERFORMANCE

Compte rendu des benchmarks fait avec JMH sur les
déclarations de performanceClaims.pdf

YANN MARTIN
D'ESCRIBENNE
N°ETUDIANT
21713752



Java performance

TABLE DES MATIERES

1.1	Standard code optimizations, Claim 1	2
1.1	Standard code optimizations, Claim 2.....	3
1.1	Standard code optimizations, Claim 3.....	4
1.1	Standard code optimizations, Claim 4.....	5
1.1	Standard code optimizations, Claim 5.....	6
1.1	Standard code optimizations, Claim 6.....	7
1.2	Fields and variables.....	8
1.3	String manipulation.....	9
1.4	Storing tables of constants in arrays.....	10
1.5	Methods, Claim 1	11
1.5	Methods, Claim 2	12
1.6	Sorting and searching, Claim 1	13
1.6	Sorting and searching, Claim 2	14
1.7	Exceptions.....	15
1.8	Collection classes, Claim 1	16
1.8	Collection classes, Claim 2	17
1.8	Collection classes, Claim 3	18
1.8	Collection classes, Claim 4	19
1.9	Input and output	20
1.11	Bulk array operations.....	21
1.13	Reflection.....	22



1.1 STANDARD CODE OPTIMIZATIONS, CLAIM 1

- Move loop-invariant computations out of loops. For example, avoid repeatedly computing the loop bound in a `for`-loop, like this:

```
for (int i=0; i<size()*2; i++) { ... }
```

Instead, compute the loop bound only once and bind it to a local variable, like this:

```
for (int i=0, stop=size()*2; i<stop; i++) { ... }
```

RESULTATS :

Benchmark	Mode	Cnt	Score	Error	Units
MyBenchmark.test1	avgt	5	5,500 ±	0,079	ms/op
MyBenchmark.test2	avgt	5	5,489 ±	0,030	ms/op
MyBenchmark.test3	avgt	5	≈ 10 ⁻⁶		ms/op

Test1: for normal.

Test2: for optimisé.

Test3: Pas de test.

CONCLUSION:

Pas de différence entre les deux boucles `for`.

1.1 STANDARD CODE OPTIMIZATIONS, CLAIM 2

- Do not compute the same subexpression twice:

```
if (birds.elementAt(i).isGrower()) ...  
if (birds.elementAt(i).isPullet()) ...
```

Instead, compute the subexpression once, bind the result to a variable, and reuse it:

```
Bird bird = birds.elementAt(i);  
if (bird.isGrower()) ...  
if (bird.isPullet()) ...
```

RESULTATS :

Benchmark	Mode	Cnt	Score	Error	Units
MyBenchmark.test1	avgt	5	32,323 ±	4,073	ms/op
MyBenchmark.test2	avgt	5	31,224 ±	6,087	ms/op
MyBenchmark.test3	avgt	5	≈ 10 ⁻⁶		ms/op

Test1: if normal.


Test2: if optimisé.

Test3: Pas de test.

CONCLUSION :

Pas de différence entre les deux if (j'ai utilisé une ArrayList). Notez qu'avec une LinkedList cette optimisation à plus d'impact avec un facteur ~2:

Benchmark	Mode	Cnt	Score	Error	Units
MyBenchmark.test1	avgt	5	12,433 ±	1,422	ms/op
MyBenchmark.test2	avgt	5	6,402 ±	0,428	ms/op
MyBenchmark.test3	avgt	5	≈ 10 ⁻⁶		ms/op



1.1 STANDARD CODE OPTIMIZATIONS, CLAIM 3

- Every array access requires an index check, so it is worth-while to reduce the number of array accesses. Moreover, usually the Java compiler cannot automatically optimize indexing into multidimensional arrays. For instance, every iteration of the inner (*j*) loop below recomputes the indexing `rowsum[i]` as well as the indexing `arr[i]` into the first dimension of `arr`:

```
double[] rowsum = new double[n];
for (int i=0; i<n; i++)
    for (int j=0; j<m; j++)
        rowsum[i] += arr[i][j];
```

Instead, compute these indexings only once for each iteration of the outer loop:

```
double[] rowsum = new double[n];
for (int i=0; i<n; i++) {
    double[] arri = arr[i];
    double sum = 0.0;
    for (int j=0; j<m; j++)
        sum += arri[j];
    rowsum[i] = sum;
}
```

RESULTATS :

Benchmark	Mode	Cnt	Score	Error	Units
MyBenchmark.test1	avgt	5	26,411 ±	5,339	ms/op
MyBenchmark.test2	avgt	5	25,753 ±	2,164	ms/op
MyBenchmark.test3	avgt	5	≈ 10 ⁻⁶		ms/op

Test1: Somme des matrices normal.

Test2: Somme de la matrice optimisée.

Test3: Pas de test.

CONCLUSION:

Pas de différence entre les deux sommes de matrice.



1.1 STANDARD CODE OPTIMIZATIONS, CLAIM 4

- Declare constant fields as `final static` so that the compiler can inline them and precompute constant expressions.
- Declare constant variables as `final` so that the compiler can inline them and precompute constant expressions.

RESULTATS :

Benchmark	Mode	Cnt	Score	Error	Units
MyBenchmark.test1	avgt	5	4,609	± 0,368	ms/op
MyBenchmark.test2	avgt	5	4,695	± 0,290	ms/op
MyBenchmark.test3	avgt	5	4,660	± 0,425	ms/op

Test1: Récupération du champ normal.

Test2: Récupération du champ static.

Test3: Récupération du champ static final.

CONCLUSION:

Pas de différence majeure entre la récupération de constantes.

1.1 STANDARD CODE OPTIMIZATIONS, CLAIM 5

- Replace a long if-else-if chain by a switch if possible; this is much faster.

RÉSULTATS :

Benchmark	Mode	Cnt	Score	Error	Units
MyBenchmark.test1	avgt	5	19,541 ±	0,476	ms/op
MyBenchmark.test2	avgt	5	14,257 ±	0,258	ms/op
MyBenchmark.test3	avgt	5	≈ 10 ⁻⁶		ms/op

Test1: If-else.

Test2: Switch.

Test3: Pas de test.

CONCLUSION:

On constate que le switch est plus rapide que le if-else (facteur ~1.3).

1.1 STANDARD CODE OPTIMIZATIONS, CLAIM 6

- Nothing (except obscurity) is achieved by using ‘clever’ C idioms such as performing the entire computation of a while-loop in the loop condition:

```
int year = 0;
double sum = 200.0;
double[] balance = new double[100];
while ((balance[year++] = sum *= 1.05) < 1000.0);
```

RESULTATS :

Benchmark	Mode	Cnt	Score	Error	Units
MyBenchmark.test1	avgt	5	0,323 ±	0,023	ms/op
MyBenchmark.test2	avgt	5	0,324 ±	0,018	ms/op
MyBenchmark.test3	avgt	5	$\approx 10^{-6}$		ms/op

Test1: while classique.

Test2: while façon C.

Test3: Pas de test.

CONCLUSION:

Pas de différence entre les deux while.

1.2 FIELDS AND VARIABLES

- Access to local variables and parameters in a method is much faster than access to static or instance fields. For a field accessed in a loop, it may be worthwhile to copy the field's value to a local variable before the loop, and refer only to the local variable inside the loop.
- There is no runtime overhead for declaring variables inside nested blocks or loops in a method. It usually improves clarity to declare variables as locally as possible (with as small a scope as possible), and this may even help the compiler improve your program.

RESULTATS :

Benchmark	Mode	Cnt	Score	Error	Units
MyBenchmark.test1	avgt	5	0,161 ± 0,003		ms/op
MyBenchmark.test2	avgt	5	0,161 ± 0,002		ms/op
MyBenchmark.test3	avgt	5	0,161 ± 0,003		ms/op

Test1: Somme avec variable statique.

Test2: Somme avec variable.

Test3: Somme avec variable locale.

CONCLUSION:

Pas de différence entre les sommes.

1.3 STRING MANIPULATION

- Do not build strings by repeated string concatenation. The loop below takes time quadratic in the number of iterations and most likely causes heap fragmentation as well (see Section 2):

```
String s = "";
for (int i=0; i<n; i++) {
    s += "#" + i;
}
```

Instead, use a `StringBuilder` object and its `append` method. This takes time linear in the number of iterations, and may be several orders of magnitude faster:

```
StringBuilder sbuf = new StringBuilder();
for (int i=0; i<n; i++) {
    sbuf.append("#").append(i);
}
String s = sbuf.toString();
```

RESULTATS :

Benchmark	Mode	Cnt	Score	Error	Units
MyBenchmark.test1	avgt	5	5,314 ±	0,427	ms/op
MyBenchmark.test2	avgt	5	0,055 ±	0,001	ms/op
MyBenchmark.test3	avgt	5	≈ 10 ⁻⁶		ms/op

Test1: Concaténation classique.

Test2: Concaténation avec un String Builder.

Test3: Pas de test

CONCLUSION:

On constate une énorme optimisation de la concaténation avec un String Builder (facteur 100).

1.4 STORING TABLES OF CONSTANTS IN ARRAYS

- Declaring an initialized array variable inside a method causes a new array to be allocated at every execution of the method:

```
public static int monthdays(int y, int m) {  
    int[] monthlengths =  
        { 31, 28, 31, 30, 31, 30, 31, 31, 30, 31, 30, 31 };  
    return m == 2 && leapyear(y) ? 29 : monthlengths[m-1];  
}
```

Instead, an initialized array variable or similar table should be declared and allocated once and for all as a final static field in the enclosing class:

```
private final static int[] monthlengths =  
    { 31, 28, 31, 30, 31, 30, 31, 31, 30, 31, 30, 31 };  
  
public static int monthdays(int y, int m) {  
    return m == 2 && leapyear(y) ? 29 : monthlengths[m-1];  
}
```

RESULTATS :

Benchmark	Mode	Cnt	Score	Error	Units
MyBenchmark.test1	avgt	5	1132,733 ± 126,014		ms/op
MyBenchmark.test2	avgt	5	64,663 ± 1,403		ms/op
MyBenchmark.test3	avgt	5	≈ 10 ⁻⁶		ms/op

Test1: Liste de constante interne à la méthode.

Test2: Liste de constante statique.

Test3: Pas de test

CONCLUSION:

On constate une énorme optimisation avec une liste de constante non recalculée (facteur ~17).

1.5 METHODS, CLAIM 1

- Declaring a method as `private`, `final`, or `static` makes calls to it faster. Of course, you should only do this when it makes sense in the application.
- For instance, often an accessor method such as `getSize` can reasonably be made `final` in a class, when there would be no point in overriding it in a subclass:

```
class Foo {  
    private int size;  
    ...  
    public final int getSize() {  
        return size;  
    }  
}
```

This can make a call `o.getSize()` just as fast as a direct access to a public field `o.size`. There need not be any performance penalty for proper encapsulation (making fields `private`).

RÉSULTATS :

Benchmark	Mode	Cnt	Score	Error	Units
MyBenchmark.test1	avgt	5	13,827	± 1,063	ms/op
MyBenchmark.test2	avgt	5	13,370	± 0,252	ms/op
MyBenchmark.test3	avgt	5	13,352	± 0,174	ms/op

Test1: get classique.

Test2: get final.

Test3: récupération du champ directement.

CONCLUSION:

On ne constate qu'une petite différence entre la méthode de get classique et la méthode de get final. Mais ce n'est pas assez notable.

1.5 METHODS, CLAIM 2

- Virtual method calls (to instance methods) are fast and should be used instead of `instanceof` tests and casts.

RESULTATS :

Benchmark	Mode	Cnt	Score	Error	Units
MyBenchmark.test1	avgt	5	3,011 ±	1,572	ms/op
MyBenchmark.test2	avgt	5	3,022 ±	1,410	ms/op
MyBenchmark.test3	avgt	5	$\approx 10^{-6}$		ms/op

Test1: appel de la méthode après instance of et cast.

Test2: appel de la méthode virtuelle.

Test3: Pas de test.

CONCLUSION:

On ne constate pas de différence entre l'appel d'une méthode après instance of puis cast et une méthode virtuelle. Ce qui est logique car il est dit plus bas dans le document :

- For example, a test (`x instanceof C`) conditionally followed by a cast (`(C) x`) may be optimized by a JVM so that at most one test is performed. Hence it is not worth the trouble to rewrite your program to avoid either the `instanceof` test or the cast.

1.6 SORTING AND SEARCHING, CLAIM 1

- Never use selection sort, bubblesort or insertion sort, except on very short arrays or lists. Use heapsort (for arrays) or mergesort (for doubly linked lists) or quicksort (for arrays; but you must make a good choice of pivot element).
- Even better, use the built-in sorting routines, which are guaranteed to be fast: $O(n \log(n))$ time for n elements, and sometimes faster if the data are nearly sorted already:

A

RÉSULTATS :

Benchmark	Mode	Cnt	Score	Error	Units
SortBenchmark.longBubbleSort	avgt	5	15,709 ±	0,309	ms/op
SortBenchmark.longHeapSort	avgt	5	0,391 ±	0,007	ms/op
SortBenchmark.longInsertionSort	avgt	5	1,613 ±	0,112	ms/op
SortBenchmark.longJavaSort	avgt	5	0,241 ±	0,003	ms/op
SortBenchmark.longMergeSort	avgt	5	0,407 ±	0,036	ms/op
SortBenchmark.longQuickSort	avgt	5	0,267 ±	0,004	ms/op
SortBenchmark.longSelectionSort	avgt	5	4,602 ±	0,057	ms/op
SortBenchmark.shortBubbleSort	avgt	5	0,058 ±	0,001	ms/op
SortBenchmark.shortHeapSort	avgt	5	0,013 ±	0,001	ms/op
SortBenchmark.shortInsertionSort	avgt	5	0,008 ±	0,001	ms/op
SortBenchmark.shortJavaSort	avgt	5	0,008 ±	0,001	ms/op
SortBenchmark.shortMergeSort	avgt	5	0,014 ±	0,001	ms/op
SortBenchmark.shortQuickSort	avgt	5	0,010 ±	0,001	ms/op
SortBenchmark.shortSelectionSort	avgt	5	0,023 ±	0,001	ms/op

CONCLUSION:

Dans des listes courtes, le tri par insertion est bon mais pas sur des longues listes. Le tri de Java quant à lui est bon dans toutes les circonstances. Viens ensuite le quickSort. Le tri à bulle est le moins bon de tous.

1.6 SORTING AND SEARCHING, CLAIM 2

- Avoid linear search in arrays and lists, except when you know that they are very short. If your program needs to look up something frequently, use one of these approaches:
 - *Binary search on sorted data:*
For arrays, use `java.util.Arrays.binarySearch`. The array must be sorted, as if by `java.util.Arrays.sort`. There are overloaded versions for all primitive types and for objects.
For `ArrayList<T>`, use `java.util.Collections.binarySearch`. The array list must be sorted, as if by `java.util.Collections.sort`.
If you need also to insert or remove elements from the set or map, use one of the approaches below instead.
 - *Hashing:* Use `HashSet<T>` or `HashMap<K, V>` from package `java.util` if your key objects have a good hash function `hashCode`. This is the case for `String` and the wrapper classes `Integer`, `Double`, ..., for the primitive types.
 - *Binary search trees:* Use `TreeSet<T>` or `TreeMap<K, V>` from package `java.util` if your key objects have a good comparison function `compareTo`. This is the case for `String` and the wrapper classes `Integer`, `Double`, ..., for the primitive types.

RÉSULTATS :

Benchmark	Mode	Cnt	Score	Error	Units
SearchBenchmark.BinarySearchArray	avgt	5	0,966 ±	0,072	ms/op
SearchBenchmark.HashMapSearch	avgt	5	0,026 ±	0,001	ms/op
SearchBenchmark.TreeMapSearch	avgt	5	0,136 ±	0,017	ms/op
SearchBenchmark.regularSearchArray	avgt	5	1528,685 ±	85,134	ms/op

CONCLUSION:

La recherche d'élément de base est absolument à éviter. Si les objets à rechercher répondent aux critères cités plus haut. Il vaut mieux privilégier une `HashMap` ou un `TreeMap`. Et dans le cas contraire une recherche dichotomique.

1.7 EXCEPTIONS

- The creation `new Exception(...)` of an exception object builds a stack trace, which is costly in time and space, and especially so in deeply recursive method calls. The creation of an object of class `Exception` or a subclass of `Exception` may be between 30 and 100 times slower than creation of an ordinary object. On the other hand, using a `try-catch` block or throwing an exception is fast.
- You can prevent the generation of this stack trace by overriding method `fillInStackTrace` in subclasses of `Exception`, as shown below. This makes creation exception instances roughly 10 times faster.

```
class MyException extends Exception {  
    public Throwable fillInStackTrace() {  
        return this;  
    }  
}
```

- Thus you should create an exception object only if you actually intend to throw it. Also, do not use exceptions to implement control flow (end of data, termination of loops); use exceptions only to signal errors and exceptional circumstances (file not found, illegal input format, and so on). If your program *does* need to throw exceptions very frequently, reuse a single pre-created exception object.

RÉSULTATS :

Benchmark	Mode	Cnt	Score	Error	Units
MyBenchmark.test1	avgt	5	168,330 ± 12,940		ms/op
MyBenchmark.test2	avgt	5	1,221 ± 0,182		ms/op
MyBenchmark.test3	avgt	5	0,007 ± 0,001		ms/op

Test1: Exception classique.

Test2: Exception modifiée.

Test3: Exception modifiée créer une seule fois avant la boucle.

CONCLUSION:

Il est nécessaire d'override la méthode citée plus haut lors d'une création d'une exception manuelle qui ne nécessite pas la trace du stack. De plus si elle sera appelée souvent, il faut toujours utiliser la même exception créer une et une seule fois.

1.8 COLLECTION CLASSES, CLAIM 1

- Beware that indexing into a `LinkedList<T>` is not a constant-time operation. Hence the loop below takes time quadratic in the size of the list `lst` if `lst` is a `LinkedList<T>`, and should not be used:

```
int size = lst.size();
for (int i=0; i<size; i++)
    System.out.println(lst.get(i));
```

Instead, use the enhanced `for` statement to iterate over the elements. It implicitly uses the collection's iterator, so the traversal takes linear time:

```
for (T x : lst)
    System.out.println(x);
```

RÉSULTATS :

Benchmark	Mode	Cnt	Score	Error	Units
MyBenchmark.test1	avgt	5	7,079 ±	0,677	ms/op
MyBenchmark.test2	avgt	5	0,017 ±	0,002	ms/op
MyBenchmark.test3	avgt	5	≈ 10 ⁻⁶		ms/op

Test1: for classique avec `get(i)`

Test2: `foreach` dans la `LinkedList`

Test3: Pas de test.

CONCLUSION:

Le `foreach` est beaucoup plus rapide avec un facteur 700.

1.8 COLLECTION CLASSES, CLAIM 2

- Repeated calls to `add(int i, T x)` or `remove(int i)` on `LinkedList<T>` should be avoided, except when `i` is at the *end or beginning* of the linked list; both perform a linear traversal to get to the `i`'th element.

RÉSULTATS :

Benchmark	Mode	Cnt	Score	Error	Units
MyBenchmark.test1	avgt	5	0,085 ±	0,070	ms/op
MyBenchmark.test2	avgt	5	25,463 ±	3,318	ms/op
MyBenchmark.test3	avgt	5	$\approx 10^{-6}$		ms/op

Test1: add et remove au début de la LinkedList

Test2: add et remove au milieu de la LinkedList

Test3: Pas de test.

CONCLUSION:

Il est beaucoup plus rapide d'ajouter à la fin ou au début d'une LinkedList plutôt qu'au milieu. (Facteur ~300)

1.8 COLLECTION CLASSES, CLAIM 3

- Repeated calls to `add(int i, T x)` or `remove(int i)` on `ArrayList<T>` should be avoided, except when `i` is at the *end* of the `ArrayList<T>`; it needs to move all elements after `i`.

RÉSULTATS :

Benchmark	Mode	Cnt	Score	Error	Units
MyBenchmark.test1	avgt	5	2,562 ±	0,255	ms/op
MyBenchmark.test2	avgt	5	0,107 ±	0,027	ms/op
MyBenchmark.test3	avgt	5	$\approx 10^{-6}$		ms/op


Test1: add et remove au début de l'ArrayList

Test2: add et remove à la fin de l'ArrayList

Test3: Pas de test.

CONCLUSION:

Il est beaucoup plus rapide d'ajouter à la fin d'une l'ArrayList plutôt qu'au début pour les raisons citées plus haut. (Facteur ~25)



1.8 COLLECTION CLASSES, CLAIM 4

- The collection classes can store only reference type data, so a value of primitive type such as `int`, `double`, ... must be wrapped as an `Integer`, `Double`, ... object before it can be stored or used as a key in a collection. This takes time and space and may be unacceptable in memory-constrained embedded applications. Note that strings and arrays are reference type data and need not be wrapped.

If you need to use collections that have primitive type elements or keys, consider using the Trove library, which provides special-case collections such as hash set of `int` and so on. As a result it is faster and uses less memory than the general Java collection classes. Trove can be found at <http://trove4j.sourceforge.net/>.

RÉSULTATS :

Benchmark	Mode	Cnt	Score	Error	Units
MyBenchmark.test1	avgt	5	41,046 ± 16,331		ms/op
MyBenchmark.test2	avgt	5	337,214 ± 70,208		ms/op
MyBenchmark.test3	avgt	5	≈ 10 ⁻⁶		ms/op

Test1: add avec une arraylist de trove (int)

Test2: add avec une arraylist de java (Integer)

Test3: Pas de test.

CONCLUSION:

On est beaucoup plus rapide en utilisant les types primitifs comme dit plus haut (Facteur ~8)

1.9 INPUT AND OUTPUT

- Using buffered input and output (`BufferedReader`, `BufferedWriter`, `BufferedInputStream`, `BufferedOutputStream` from package `java.io`) can speed up input/output by a factor of 20.
- Using the compressed streams `ZipInputStream` and `ZipOutputStream` or `GZIPInputStream` and `GZIPOutputStream` from package `java.util.zip` may speed up the input and output of verbose data formats such as XML. Compression and decompression takes CPU time, but the compressed data may be so much smaller than the uncompressed data that it saves time anyway, because less data must be read from disk or network. Also, it saves space on disk.

RÉSULTATS :

Benchmark	Mode	Cnt	Score	Error	Units
MyBenchmark.test1	avgt	5	508,403 ± 13,451		ms/op
MyBenchmark.test2	avgt	5	225,387 ± 4,552		ms/op
MyBenchmark.test3	avgt	5	473,749 ± 10,189		ms/op

Test1: avec un Reader et Writer classique

Test2: avec un `BufferedReader` et `BufferedWriter`

Test3: avec un `BufferedReader` et `BufferedWriter` et compression en Zip.

CONCLUSION:


L'utilisation des `BufferedReader` et `BufferedWriter` est beaucoup plus rapide que le reste (Facteur ~2), mais la compression en zip réduit la taille du fichier drastiquement.

text
text.zip

18/10/2020 12:02
18/10/2020 12:02

Fichier
Archive WinRAR ZIP

8 192 Ko
50 Ko



1.11 BULK ARRAY OPERATIONS

There are special methods for performing bulk operations on arrays. They are usually much faster than equivalent `for` loops, in part because they need to perform only a single bounds check.

- `static void java.lang.System.arraycopy(src, si, dst, di, n)` copies elements from array segment `src[si..si+n-1]` to array segment `dst[di..di+n-1]`.
- `static bool java.util.Arrays.equals(arr1, arr2)` returns true if the arrays `arr1` and `arr2` have the same length and their elements are pairwise equal. There are overloads of this method for arguments of type `boolean[], byte[], char[], double[], float[], int[], long[], Object[]` and `short[]`.
- `static void java.util.Arrays.fill(arr, x)` sets all elements of array `arr` to `x`. This method has the same overloads as `Arrays.equals`.
- `static void java.util.Arrays.fill(arr, i, j, x)` sets elements `arr[i..j-1]` to `x`. This method has the same overloads as `Arrays.equals`.

RÉSULTATS :

Benchmark	Mode	Cnt	Score	Error	Units
ArrayMethodBenchmark.ArrayCopy	avgt	5	13,518	± 2,859	ms/op
ArrayMethodBenchmark.ArrayEqual	avgt	5	18,933	± 1,102	ms/op
ArrayMethodBenchmark.ArrayFill	avgt	5	10,015	± 0,633	ms/op
ArrayMethodBenchmark.regularCopy	avgt	5	12,821	± 1,146	ms/op
ArrayMethodBenchmark.regularEqual	avgt	5	16,001	± 0,796	ms/op
ArrayMethodBenchmark.regularFill	avgt	5	9,893	± 0,982	ms/op

CONCLUSION:

Il n'y a pas de différence majeure entre les méthodes faites à la main et celle natives de `java.util.Arrays`

1.13 REFLECTION

- A reflective method call, reflective field access, and reflective object creation (using package `java.lang.reflect`) are far slower than ordinary method call, field access, and object creation.
- Access checks may further slow down such reflective calls; some of this cost may be avoided by declaring the class of the called method to be `public`. This has been seen to speed up reflective calls by a factor of 8.

RÉSULTATS :

Benchmark	Mode	Cnt	Score	Error	Units
MyBenchmark.test1	avgt	5	3,001 ±	0,305	ms/op
MyBenchmark.test2	avgt	5	55,099 ±	6,146	ms/op
MyBenchmark.test3	avgt	5	61,364 ±	10,451	ms/op

Test1: appel de la méthode normalement.

Test2: appel de la méthode par réflexion.

Test3: appel de la méthode par réflexion avec un check de l'accessibilité avant.

CONCLUSION:

En passant par la réflexion, on a un facteur ~ 18 . Mais en regardant l'accessibilité avant on ne perd pas énormément de temps contrairement à ce qui est dit au-dessus.